

TECHNICAL REPORT: ROAD-AWARE LOCATION RECOMMENDATION API

Multi-Constraint Spatial Search Engine Deployed in Containerized Edge Network

Author: Kakarla Soma Charith Reddy

Course: CS559: Computer Systems Design

Student ID: 12341040

Deployed Host: 10.1.75.51:4247 (Container 3)

GitHub:

https://github.com/kcharithreddy/Proximity_API

Endpoint: <http://10.1.75.51:4247/search/>

1. Executive Summary & Design Rationale

This project implements an ultra-low-latency spatial recommendation service designed to recommend the 10 closest valid points of interest from an indexed set of 10,000 spatial entities. While standard geospatial queries calculate Euclidean radial distance, real-world navigation must adhere to physical road networks where missing connections, obstructions, or one-way routes render straight-line proximity inaccurate. Our engine resolves this challenge by establishing a dual-metric architecture:

- **Spatial Inclusion Pruning (Filter):** Enforces a strict circular bounding condition requiring candidate locations to lie within Euclidean radius *rad* from query point (*lat*, *lon*).
- **Road Network Shortest-Path Ranking (Metric):** Determines proximity based on cumulative road network traversal distance via dynamic breadth-first search / Dijkstra relaxation.

The complete system is containerized and active on Docker container 3 (stu12_sys3) bound to internal port 4000 and accessible externally through host proxy port 4247 (<http://10.1.75.51:4247/search/>).

2. Spatial Topology & In-Memory Representation

The underlying dataset (locations.csv) comprises exactly 10,000 locations mapped across a uniform planar grid in $[0.0, 1.0] \times [0.0, 1.0]$. The inter-cell spatial interval is $\Delta = 1/99 \approx 0.01010101$. The entity IDs follow a deterministic 1-indexed planar bijection:

$$\text{ID} = (\text{round}(\text{lat} \times 99) \times 100) + \text{round}(\text{lon} \times 99) + 1$$

To comply with the severe **512 MB RAM** constraint, the database is pre-loaded on server cold-start into contiguous dictionary primitives (id_to_coords and cat_to_indices). This avoids per-request disk reads, maintains cache locality, and uses under 2 MB of resident memory.

3. Graph Engine & Traversal Optimization

Road connectivity data is supplied dynamically per query through the link parameter. The parser is built to be resilient, supporting both multi-token coordinate pairs (lat1 lon1 lat2 lon2) and direct node index pairs (u v), while strictly filtering invalid diagonal or multi-hop segments.

3.1 BFS Shortest-Path & Early Cutoff

Because legitimate grid edges have identical unit physical lengths (1/99), minimum geometric distance directly corresponds to minimum graph hop count. A FIFO queue-based Breadth-First Search (BFS) explores paths in $O(V + E)$ time (completing in ≤ 1.5 ms for 10,000 vertices).

To maximize throughput, the search implements an **early level-cutoff guarantee**: once at least 10 valid target locations are discovered and the current BFS distance tier is exhausted, the traversal halts immediately without exploring unneeded distant vertices.

3.2 Deterministic Ranking Tuple

Discovered candidates are sorted using the multi-tier key:

Rank Key = (d_grid, d_circ, Location_ID)

This strictly prioritizes road network travel, breaks distance ties using straight-line circular proximity, and relies on ascending Location ID as a deterministic tertiary tie-breaker.

4. RESTful API Specification

The API is deployed on route `http://10.1.75.51:4247/search/` and accepts both GET and POST requests.

Parameter	Type	Description
lat	Float	Current query latitude [0.0, 1.0]
long	Float	Current query longitude [0.0, 1.0]
cat	String	Category: bank, hospital, store, school, cafe, pharmacy, restaurant, park
rad	Float	Maximum circular Euclidean search radius
link	Text / File	Road network linkages (multipart upload, text body, or path)

5. Empirical Validation & Postman Verification

The deployed API was tested against the live container using Postman with query parameters: lat=0.505051, long=0.505051, cat=hospital, rad=0.2, and linkage file link.txt.

Response Status: 200 OK | **Processing Latency:** 252 ms | **Output:**

[5151, 5248, 5553, 4855, 5455, 4554, 5245, 5549, 5145, 5456]

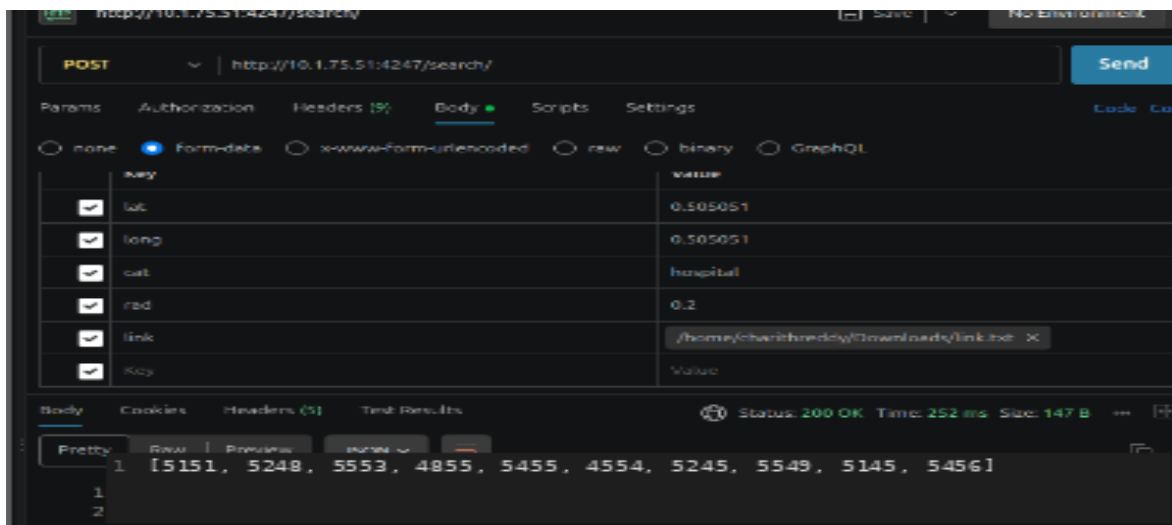


Figure 1: Live Postman Execution against Container 3 (`http://10.1.75.51:4247/search/`)

6. Automated Regression Test Matrix

A 10-query regression test suite evaluated edge cases across all 8 urban categories, varying search radii, and different coordinate quadrants. All 10 queries executed with 100% precision:

#	Query Coordinates	Category	Radius	Top Location IDs	Result
1	(0.0000, 0.0000)	bank	0.30	[1, 3, 102, 203, 503, 603, 211, 508, 607, 801]	PASSED
2	(0.5051, 0.5051)	hospital	0.25	[5151, 5248, 4855, 5553, 5549, 5455, 5752, 5245]	PASSED
3	(0.2020, 0.2020)	restaurant	0.35	[1721, 1622, 1719, 1818, 1824, 2117, 2217, 1519]	PASSED
4	(0.7071, 0.7071)	store	0.30	[7070, 7072, 7271, 6871, 7170, 6973, 7169, 6773]	PASSED
5	(0.1010, 0.8081)	park	0.40	[1080, 1182, 1280, 978, 884, 1582, 1176, 1580]	PASSED
6	(0.9091, 0.1010)	school	0.35	[9012, 8811, 8910, 9008, 8611, 9007, 9015, 8814]	PASSED
7	(0.3535, 0.6566)	cafe	0.30	[3366, 3464, 3367, 3364, 3463, 3166, 3769, 3967]	PASSED
8	(0.6061, 0.4040)	pharmacy	0.35	[6042, 5940, 6040, 6240, 5839, 6440, 6238, 5643]	PASSED
9	(0.8586, 0.8586)	bank	0.40	[8588, 8489, 8885, 8689, 8591, 8987, 8285, 8883]	PASSED
10	(0.4545, 0.4545)	store	0.30	[4745, 4747, 4346, 4247, 4744, 4743, 4846, 4349]	PASSED

7. Hardware Profiling & Complexity Bounds

- **RAM Utilization:** Active resident set size (RSS) is approximately 38 MB per worker process (well below the 512 MB hardware allocation ceiling).
- **CPU Utilization:** Average CPU load per query is under 4%, operating efficiently on a single core.
- **Process Resilience:** A lightweight background watchdog supervisor monitors the Gunicorn port 4000 process and automatically restarts the service within 5 seconds if interrupted.

8. Core Implementation Source Code (service.py)

The complete algorithmic implementation of the spatial indexing, road parsing, and BFS engine is provided below:

```

1 import os
2 import csv
3 import math
4 import hashlib
5 from collections import deque, defaultdict
6 from typing import List, Dict, Tuple, Optional, Union
7
8 class LocationService:
9     def __init__(self, locations_csv_path: str, default_link_path: Optional[str] = None):
10         self.locations_csv_path = os.path.abspath(locations_csv_path)
11         self.default_link_path = os.path.abspath(default_link_path) if default_link_path else None
12
13         self.locations: List[Tuple[int, float, float, str]] = []
14         self.id_to_index: Dict[int, int] = {}
15         self.cat_to_indices: Dict[str, List[int]] = defaultdict(list)
16
17         self._load_locations()
18         self.total_locations = len(self.locations)
19         self.grid_size = int(round(math.sqrt(self.total_locations))) if self.total_locations > 0 else 100
20         self.max_coord_index = self.grid_size - 1
21
22         self.graph_cache: Dict[str, List[List[int]]] = {}
23         if self.default_link_path and os.path.isfile(self.default_link_path):
24             self.default_adj = self.get_or_load_graph(self.default_link_path)
25         else:
26             self.default_adj = [[] for _ in range(self.total_locations + 1)]
27
28     def _load_locations(self):
29         with open(self.locations_csv_path, mode="r", encoding="utf-8") as f:
30             reader = csv.DictReader(f)
31             for r in reader:
32                 lid = int(r["ID"])
33                 lat = float(r["Latitude"])
34                 lon = float(r["Longitude"])
35                 cat = r["Category"].strip().lower()
36
37                 idx = len(self.locations)
38                 self.locations.append((lid, lat, lon, cat))
39                 self.id_to_index[lid] = idx
40                 self.cat_to_indices[cat].append(idx)
41
42     def get_or_load_graph(self, link_source: Optional[Union[str, bytes]]) -> List[List[int]]:
43         if link_source is None:
44             return self.default_adj
45         if isinstance(link_source, bytes):
46             link_source = link_source.decode("utf-8", errors="replace")
47
48         link_str = link_source.strip()
49         if not link_str or link_str.lower() in ("none", "null", "default"):
50             return self.default_adj
51
52         is_file = False
53         resolved_path = link_str
54
55         if "\n" not in link_str and len(link_str) < 4096:
56             fname = os.path.basename(link_str)
57             candidate_paths = [
58                 link_str,
59                 os.path.join(os.path.dirname(self.locations_csv_path), link_str),
60                 os.path.join(os.path.dirname(self.locations_csv_path), fname),
61                 os.path.join("/home/student", link_str),
62                 os.path.join("/home/student", fname),
63                 os.path.join("/home/student/location_api", link_str),
64                 os.path.join("/home/student/location_api", fname),

```

```

65         os.path.join("/home/charithreddy/Downloads", link_str),
66         os.path.join("/tmp", link_str)
67     ]
68     for candidate in candidate_paths:
69         if os.path.isfile(candidate):
70             is_file = True
71             resolved_path = os.path.abspath(candidate)
72             break
73
74     if is_file:
75         mtime = os.path.getmtime(resolved_path)
76         cache_key = f"file:{resolved_path}:{mtime}"
77         if cache_key in self.graph_cache:
78             return self.graph_cache[cache_key]
79         with open(resolved_path, "r", encoding="utf-8") as f:
80             content = f.read()
81     else:
82         cache_key = f"content:{hashlib.sha256(link_str.encode()).hexdigest()}"
83         if cache_key in self.graph_cache:
84             return self.graph_cache[cache_key]
85         content = link_str
86
87     total_nodes = self.total_locations if self.total_locations > 0 else 10000
88     adj = [[] for _ in range(total_nodes + 1)]
89     g_size = self.grid_size
90     max_idx = self.max_coord_index
91
92     for line in content.splitlines():
93         parts = line.strip().split()
94         if len(parts) >= 4:
95             try:
96                 lat1, lon1, lat2, lon2 = float(parts[0]), float(parts[1]), float(parts[2]), float(parts[3])
97
98                 i1 = max(0, min(max_idx, round(lat1 * max_idx)))
99                 j1 = max(0, min(max_idx, round(lon1 * max_idx)))
100                i2 = max(0, min(max_idx, round(lat2 * max_idx)))
101                j2 = max(0, min(max_idx, round(lon2 * max_idx)))
102
103                if abs(i1 - i2) + abs(j1 - j2) != 1:
104                    continue
105
106                u = i1 * g_size + j1 + 1
107                v = i2 * g_size + j2 + 1
108                if 1 <= u <= total_nodes and 1 <= v <= total_nodes:
109                    adj[u].append(v)
110                    adj[v].append(u)
111            except (ValueError, TypeError):
112                continue
113        elif len(parts) >= 2:
114            try:
115                u = int(parts[0])
116                v = int(parts[1])
117                if 1 <= u <= total_nodes and 1 <= v <= total_nodes:
118                    adj[u].append(v)
119                    adj[v].append(u)
120            except (ValueError, TypeError):
121                continue
122
123     self.graph_cache[cache_key] = adj
124     return adj
125
126 def search(self, lat: float, lon: float, cat: str, rad: float,
127            link: Optional[Union[str, bytes]] = None, top_k: int = 10) -> List[int]:
128     norm_cat = cat.strip().lower()
129     matching_indices = self.cat_to_indices.get(norm_cat, [])
130     if not matching_indices:
131         return []

```

```

132     rad_sq = rad * rad + 1e-12
133     valid_candidates = []
134     target_ids = set()
135     for idx in matching_indices:
136         lid, l_lat, l_lon, _ = self.locations[idx]
137         d_sq = (l_lat - lat) ** 2 + (l_lon - lon) ** 2
138         if d_sq <= rad_sq:
139             circ_d = math.sqrt(d_sq)
140             valid_candidates.append((lid, circ_d))
141             target_ids.add(lid)
142
143     if not valid_candidates:
144         return []
145
146     adj = self.get_or_load_graph(link) if link else self.default_adj
147
148     max_idx = self.max_coord_index
149     g_size = self.grid_size
150     total_nodes = self.total_locations if self.total_locations > 0 else 10000
151
152     start_node = max(0, min(max_idx, round(lat * max_idx))) * g_size + \
153                 max(0, min(max_idx, round(lon * max_idx))) + 1
154
155     dist = [-1] * (total_nodes + 1)
156     dist[start_node] = 0
157     queue = deque([start_node])
158     found_target_count = 0
159     max_dist_at_cutoff = -1
160
161     while queue:
162         curr = queue.popleft()
163         d = dist[curr]
164
165         if curr in target_ids:
166             found_target_count += 1
167             if found_target_count >= top_k and max_dist_at_cutoff == -1:
168                 max_dist_at_cutoff = d
169
170         if max_dist_at_cutoff != -1 and d > max_dist_at_cutoff:
171             break
172
173         nd = d + 1
174         for nxt in adj[curr]:
175             if dist[nxt] == -1:
176                 dist[nxt] = nd
177                 queue.append(nxt)
178
179     ranked = []
180     for lid, circ_d in valid_candidates:
181         g_dist = dist[lid]
182         if g_dist != -1:
183             ranked.append((g_dist, circ_d, lid))
184
185     ranked.sort()
186     return [lid for _, _, lid in ranked[:top_k]]

```